



Growing Hypervisor 0day with Hyperseed

Daniel King @long123king
Shawn Denbow @sdenbow_

MSRC Vulnerabilities & Mitigations Team
CDG Platform Security & Vulnerability Research

whoami

- Daniel King
 - MSRC V&M
 - Keen Security Lab
 - Member of “Master of Pwn” 2016
 - MSRC Nano Server Bounty Winner
 - Ruxcon 2016
 - CodeBlue 2016
 - ZeroNights 2015
- Shawn Denbow
 - CDG PSVR
 - USAF
 - Matasano
 - RPSEC

Overview

- Motivation
- Hyper-V Architecture
- Hypervisors
- Hyperseed
- Mutations
- Findings

Motivation – Why Hypervcalls?

Importance

- Cornerstone of Hyper-V

Accessibility

- Accessible from kernel (Ring-0) mode

Simplicity

- Well-documented, good starting point

Complexity

- Diverse input & output format

Motivation – Bug Bounty

RCE w/ Exploit (Guest-to-Host-Escape)

- \$250,000 (Hypervisor/Kernel) / \$150,000 (User-mode)

RCE (Guest-to-Host-Escape)

- \$200,000 (Hypervisor/Kernel) / \$100,000 (User-mode)


Information Disclosure


- \$25,000 (Hypervisor/Kernel), \$15,000 (User-mode)

Denial of Service

- \$15,000 (Hypervisor/Kernel)

Hyper-V Architecture







<https://docs.microsoft.com/en-us/virtualization/hyper-v-on-windows/reference/hyper-v-architecture>

Nested Virtualization

- L0 Hyper-V Hypervisor
- L1 Root
- L1 Guest
- L1 Hyper-V Hypervisor
- L2 Root
- L2 Guest



Virtualization Based Security



- VBS introduces concept of trust levels
 - VTL0 – normal mode
 - VTL1 – secure (privileged) mode
- VTL1 controls which resources VTL0 can access
 - Memory, privileged CPU registers
- Hypervisor leverages hardware virtualization to enforce access permissions
 - SLAT – memory access protection for CPU
 - IOMMU – memory access protection for devices
- VTL1 is NOT a separate VM
 - VTL1 and VTL0 share memory, CPU's
 - Secure kernel coordinates with normal kernel for memory management, scheduling, etc.



Hypercalls



Hypervisor Basics

- Calling mechanism for guests to interface with the hypervisor
 - Three calling conventions
 - Slow
 - Fast
 - XMM Fast
- Hypervisor classes
 - Simple
 - Rep “repeat”
- Variable Sized Hypervisors
- Extended Hypervisor Interface
 - Call codes above 0x8000

```
typedef union _HV_HYPERCALL_INPUT
{
    struct
    {
        UINT32 CallCode : 16;
        UINT32 IsFast : 1;
        UINT32 VariableHeaderSize : 9;
        UINT32 Reserved1 : 6;
        UINT32 CountOfElements : 12;
        UINT32 Reserved2 : 4;
        UINT32 RepStartIndex : 12;
        UINT32 Reserved3 : 4;
    };
    UINT64 AsUINT64;
} HV_HYPERCALL_INPUT, *PHV_HYPERCALL_INPUT;
```

```
typedef union _HV_HYPERCALL_OUTPUT
{
    struct
    {
        UINT16 CallStatus;
        UINT16 Reserved1;
        UINT32 ElementsProcessed : 12;
        UINT32 Reserved2 : 20;
    };
    UINT64 AsUINT64;
} HV_HYPERCALL_OUTPUT, *PHV_HYPERCALL_OUTPUT;
```

How To Issue a Hypervisor Call

- Hypervisors must be invoked from the most privileged guest processor mode, CPL0 / Kernel mode
- All hypervisors should be invoked through the architecturally-defined hypervisor interface
 - `vmcall` instruction for Intel VMX
 - `vmmcall` instruction for AMD SVM
- Hypervisor Page: abstract the difference between different virtualization implementations
 - `#define HV_X64_MSR_HYPERCALL 0x40000001`

Establishing the Hypervisor Interface

- Write guest id to
HV_X64_MSR_GUEST_OS_ID
- Allocate a page of memory for
the hypervisor page
- Write to the hypervisor MSR
HV_X64_MSR_HYPERCALL
 - Set the Enable Hypervisor Page bit
 - Write the GPA of the hypervisor page

```
/*
 * Setup the hypervisor page and enable hypervisors.
 * 1. Register the guest ID
 * 2. Enable the hypervisor and register the hypervisor page
 */
guest_id = generate_guest_id(0, LINUX_VERSION_CODE, 0);
wrmsrl(HV_X64_MSR_GUEST_OS_ID, guest_id);


hypervisor_pg = __vmalloc(PAGE_SIZE, GFP_KERNEL, PAGE_KERNEL_RX);

if (hypervisor_pg == NULL) {
    wrmsrl(HV_X64_MSR_GUEST_OS_ID, 0);
    return;
}

rdmsrl(HV_X64_MSR_HYPERCALL, hypervisor_msr.as_uint64);
hypervisor_msr.enable = 1;
hypervisor_msr.guest_physical_address = vmalloc_to_pfn(hypervisor_pg);
wrmsrl(HV_X64_MSR_HYPERCALL, hypervisor_msr.as_uint64);
```

FROM: <https://github.com/LIS/lis-next>

Hypervisor Page



Calling here will issue a hypercall

Calling here will issue a secure call
from VTL0 to VTL1

Calling here will issue a normal call
from VTL1 to VTL0

HvCallVtlCall & HvCallVtlReturn

VTL 0 -> VTL 1

- nt!VslpEnterLumSecureMode
- vmcall(0x11)




VTL 1 -> VTL 0

- sk!SkCallNormalMode
- vmcall(0x12)

VTL Call & Return Buffer

- A [uint8_t]: SecureOperation
 - 2: Invoke Secure Service
- B [uint8_t]: NormalOperation
 - 0: Invoke Normal Service
- C [uint16_t]: ServiceCode
- D [uint64_t]: Parameters[12]



Restrictions Enforced by Hypervisor Handlers

Restrictions Enforced by Hypervisor Handlers

- The calling partition must possess a particular privilege

Partition Properties

- Partitions are identified by using a partition ID
 - `#define HV_PARTITION_ID_INVALID ((HV_PARTITION_ID) 0x0)`
 - `#define HV_PARTITION_ID_SELF ((HV_PARTITION_ID) -1)`
- Each partition has a set of properties
 - Read a property of a partition by invoking `HvGetPartitionProperty`
 - Modify partitions property by invoking `HvSetPartitionProperty`
 - `HvPartitionPropertyPrivilegeFlags` defines what the partition is allowed to access
- Partition privilege flags represented by a bit-field

Partition Privilege Flags

```
// Access to virtual MSRs                                // Access to hypercalls
AccessVpRunTimeReg:1;                                    CreatePartitions:1;
AccessPartitionReferenceCounter:1;                        AccessPartitionId:1;
AccessSynicRegs:1;                                       AccessMemoryPool:1;
AccessSyntheticTimerRegs:1;                             AdjustMessageBuffers:1;
AccessIntrCtrlRegs:1;                                    PostMessages:1;
AccessHypercallMsrs:1;                                  SignalEvents:1;
AccessVpIndex:1;                                         CreatePort:1;
AccessResetReg:1;                                       ConnectPort:1;
AccessStatsReg:1;                                       AccessStats:1;
AccessPartitionReferenceTsc:1;                           Reserved2:2;
AccessGuestIdleReg:1;                                   Debugging:1;
AccessFrequencyRegs:1;                                 CpuManagement:1;
AccessDebugRegs:1;                                     Reserved:1;
AccessReenlightenmentControls:1;                       Reserved:1;
Reserved1:18;                                           Reserved:1;
                                                               AccessVsm:1;
                                                               AccessVpRegisters:1;
                                                               Reserved:1;
                                                               Reserved:1;
                                                               EnableExtendedHypercalls:1;
                                                               StartVirtualProcessor:1;
                                                               Reserved3:10;
```

Partition Privilege Flags - Child

```
// Access to virtual MSRs                                // Access to hypercalls
AccessVpRunTimeReg:1;                                    CreatePartitions:1;
AccessPartitionReferenceCounter:1;                        AccessPartitionId:1;
AccessSynicRegs:1;                                       AccessMemoryPool:1;
AccessSyntheticTimerRegs:1;                             AdjustMessageBuffers:1;
AccessIntrCtrlRegs:1;                                    PostMessages:1;
AccessHypercallMsrs:1;                                  SignalEvents:1;
AccessVpIndex:1;                                         CreatePort:1;
AccessResetReg:1;                                       ConnectPort:1;
AccessStatsReg:1;                                       AccessStats:1;
AccessPartitionReferenceTsc:1;                           Reserved2:2;
AccessGuestIdleReg:1;                                   Debugging:1;
AccessFrequencyRegs:1;                                 CpuManagement:1;
AccessDebugRegs:1;                                     Reserved:1;
AccessReenlightenmentControls:1;                       Reserved:1;
Reserved1:18;                                           Reserved:1;
                                                               AccessVsm:1;
                                                               AccessVpRegisters:1;
                                                               Reserved:1;
                                                               Reserved:1;
                                                               EnableExtendedHypercalls:1;
                                                               StartVirtualProcessor:1;
                                                               Reserved3:10;
```

Partition Privilege Flags - Child

```
// Access to virtual MSRs                                // Access to hypercalls
AccessVpRunTimeReg:1;                                    CreatePartitions:1;
AccessPartitionReferenceCounter:1;                        AccessPartitionId:1;
AccessSyncRegs:1;                                       AccessMemoryPool:1;
AccessSyntheticTimerRegs:1;                             AdjustMessageBuffers:1;
AccessIntrCtrlRegs:1;                                   PostMessages:1;
AccessHypercallMsrs:1;                                 SignalEvents:1;
AccessVpIndex:1;                                       CreatePort:1;
AccessResetReg:1;                                      ConnectPort:1;
AccessStatsReg:1;                                       AccessStats:1;
AccessPartitionReferenceTsc:1;                           Reserved2:2;
AccessGuestIdleReg:1;                                  Debugging:1;
AccessFrequencyRegs:1;                                CpuManagement:1;
AccessDebugRegs:1;                                     Reserved:1;
AccessReenlightenmentControls:1;                      Reserved:1;
Reserved1:18;                                         Reserved:1;
                                                               AccessVsm:1;
                                                               AccessVpRegisters:1;
                                                               Reserved:1;
                                                               Reserved:1;
                                                               EnableExtendedHypercalls:1;
                                                               StartVirtualProcessor:1;
                                                               Reserved3:10;
```

Partition Privilege Flags - Root

```
// Access to virtual MSRs                                // Access to hypercalls
AccessVpRunTimeReg:1;                                    CreatePartitions:1;
AccessPartitionReferenceCounter:1;                        AccessPartitionId:1;
AccessSynicRegs:1;                                       AccessMemoryPool:1;
AccessSyntheticTimerRegs:1;                             AdjustMessageBuffers:1;
AccessIntrCtrl1Regs:1;                                   PostMessages:1;
AccessHypercallMsrs:1;                                  SignalEvents:1;
AccessVpIndex:1;                                         CreatePort:1;
AccessResetReg:1;                                       ConnectPort:1;
AccessStatsReg:1;                                       AccessStats:1;
AccessPartitionReferenceTsc:1;                           Reserved2:2;
AccessGuestIdleReg:1;                                   Debugging:1;
AccessFrequencyRegs:1;                                 CpuManagement:1;
AccessDebugRegs:1;                                     Reserved:1;
AccessReenlightenmentControls:1;                         Reserved:1;
Reserved1:18;                                           Reserved:1;
                                                               AccessVsm:1;
                                                               AccessVpRegisters:1;
                                                               Reserved:1;
                                                               Reserved:1;
                                                               EnableExtendedHypercalls:1;
                                                               StartVirtualProcessor:1;
                                                               Reserved3:10;
```

Partition Privilege Check

```
//  
// Check for CpuManagement permission.  
//  
status = HV_STATUS_SUCCESS;  
privileges.AsUINT64 = 0;  
privileges.CpuManagement = TRUE;  
HT_CHK(status = VmCheckCurrentPartitionPrivileges(privileges));
```

Restrictions Enforced by Hypervisor Handlers

- The calling partition must possess a particular privilege
- The partition being acted upon must be in a particular state
- The partition must be either a parent or child

Partition State Check

```
//  
// This function locates a partition with the specified partition Id and  
// references it only if its state matches one of the specified values.  
//  
HT_CHK(status = ObReferencePartitionById(InputHeader->PartitionId,  
                                         ObPartitionReferenceIfActive,  
                                         ObPartitionRightParent,  
                                         ObPartitionSecurityNormal,  
                                         &partition,  
                                         &partitionReference));
```

Restrictions Enforced by Hypervisor Handlers

- The calling partition must possess a particular privilege
- The partition being acted upon must be in a particular state
- The partition must be either a parent or child
- The partition must be the root

Root Partition Check

```
//  
// This hypercall can only be made by the highest enabled VTL in the root partition.  
//  
partition = VmGetCurrentPartition();  
if (!partition->Settings.IsRootPartition)  
{  
    HT_ERR(status = HV_STATUS_ACCESS_DENIED);  
}  
HT_CHK(status = VmCheckCurrentVtlAuthority(VmRootPartition, VmGetPtHighestEnabledVtl()));
```

Restrictions Enforced by Hypervisor Handlers

- The calling partition must possess a particular privilege
- The partition being acted upon must be in a particular state
- The partition must be the root
- The partition must be either a parent or child
- The virtual processor must be in a particular state


Virtual Processor State Check

```
//  
// VP must be explicitly suspended before restore.  
//  
if (!Info->Vp->Activity.ExplicitSuspend)  
{  
    HT_ERR(status = HV_STATUS_INVALID_VP_STATE);  
}
```


Hyperseed




Fuzzing from Guest




Fuzzing from Guest




Fuzzing from Root




Fuzzing from Root




Fuzzing from L2 Root Partition




Fuzzing from L2 Root Partition



Fuzzing from Host VTL 0



Fuzzing from Host VTL 0



Format-Aware Fuzzing

Past Fuzzing Experience

- CLFS, Deathnote of Microsoft Windows Kernel
- Intel PT & kAFL


Pros

- Effective to bypass fields constraints
- Easy to scale, from coarse-grained to fine-grained, from less hypercalls to more hypercalls.

Cons

- Assumptions may overlook some special circumstance
- Knowledge of input format needed


Knowledge Accumulation Cycle




How To Mutate?




Primitive Data Types Mutator (uint64_t)




Primitive Data Types Mutator (uint64_t)




Primitive Data Types Mutator (uint64_t)



Primitive Data Types Mutator (uint64_t)



Primitive Data Types Mutator (uint64_t)




Mutator Dispatcher

```
class CMutatorDispatcher
{
public:
    template<class T>
    static bool mutateBasicType(T& data, size_t hint = 0)
    {
        size_t bits = CPatternGenerator::Instance()->next_rand();
        size_t index = 0;

        if (is_bit_set(bits, index++)) CRandomMutator::mutateBasicType<T>(data, hint);
        if (is_bit_set(bits, index++)) CBitFlipMutator::mutateBasicType<T>(data, 1);
        if (is_bit_set(bits, index++)) CSLidingMutator::mutateBasicType<T>(data, hint);
        if (is_bit_set(bits, index++)) CBitFlipMutator::mutateBasicType<T>(data, 2);
        if (is_bit_set(bits, index++)) CSwapHalfMutator::mutateBasicType<T>(data, hint);
        if (is_bit_set(bits, index++)) CBitFlipMutator::mutateBasicType<T>(data, 4);
        if (is_bit_set(bits, index++)) CBitFlipMutator::mutateBasicType<T>(data, 8);
        if (is_bit_set(bits, index++)) CMaskMutator::mutateBasicType<T>(data, hint);

        return true;
    }
}
```

Hyperseed Components



Hyperseed.exe : issue_hypervisorcall()

HypercallProxy.sys : vmcall

CFlushVASMutation

Wrapper Interface

```
HV_STATUS  
HvFlushVirtualAddressSpace(  
    __in HV_ADDRESS_SPACE_ID      AddressSpace,  
    __in HV_FLUSH_FLAGS          Flags,  
    __in UINT64 ProcessorMask  
) ;
```


Native Interface

HvFlushVirtualAddressSpace	
Call Code = 0x0002	
► Input Parameters	
0	AddressSpace (8 bytes)
8	Flags (8 bytes)
16	ProcessorMask (8 bytes)

```
#define SELECT CPatternGenerator::Instance()->select  
#define FROM_RANGE CPatternGenerator::Instance()->from_range  
#define RAND_8() CPatternGenerator::Instance()->next_rand()  
#define RAND_16() CPatternGenerator::Instance()->next_rand_word()  
#define RAND_32() CPatternGenerator::Instance()->next_rand_dword()  
#define RAND_64() CPatternGenerator::Instance()->next_rand_qword()  
#define BM(x) CMutatorDispatcher::mutateBasicType(x)  
#define BMN0(x) CMutatorDispatcher::mutateBasicType_Nonzero(x)  
#define DM(x, y) CMutatorDispatcher::mutateData(x, y)
```

```
template<>  
bool mutateCustom<HvCallFlushVirtualAddressSpaceTrait>(CBaseMutation* Mutation,  
    HvCallFlushVirtualAddressSpaceBuffer& m_HypercallBuffer,  
    shared_ptr<CBaseFeed> accum_feed)  
{  
    PHV_INPUT_FLUSH_VIRTUAL_ADDRESS_SPACE input = m_HypercallBuffer.InputBuffer();  
  
    BM(input->Header.AddressSpace);  
    // #define HV_FLUSH_ALL_PROCESSORS (0x00000001)  
    // #define HV_FLUSH_ALL_VIRTUAL_ADDRESS_SPACES (0x00000002)  
    // #define HV_FLUSH_NON_GLOBAL_MAPPINGS_ONLY (0x00000004)  
    // #define HV_FLUSH_USE_EXTENDED_RANGE_FORMAT (0x00000008)  
    input->Header.Flags = RAND_8() & 0x07;  
    BMN0(input->Header.ProcessorMask);  
  
    accum_feed->set_input(string((char*)input, m_HypercallBuffer.input_len));  
    return true;  
}
```


Findings




Fuzzing Stats

- Initiated 1 Year Ago
- 2 Contributors
- 350+ Source Files, 16K+ Lines of Code
- 169 Hypercalls Mutated
- 16 Findings So Far
 - 8 EoP/RCE
 - 8 DoS
 - 2 Qualified for Hyper-V Bounty
 - 7 Qualified for WIP Bounty

Findings



How Much Are They Worth? (Hyper-V Bounty)



FROM: <https://www.microsoft.com/en-us/msrc/bounty-hyper-v>

Remote Code Execution

An eligible submission includes a RCE vulnerability in Microsoft Hyper-V that enables a guest virtual machine to compromise the hypervisor, escape from a guest virtual machine to the host, or escape from one guest virtual machine to another guest virtual machine.

Vulnerability Type	Tier	Proof of concept	Functioning Exploit	Report Quality	Payout range (USD)*
RCE	Tier 1	Required	Yes	High	\$250,000
			No	High	\$200,000
RCE	Tier 2	Required	Yes	High	\$150,000
			No	High	\$100,000
RCE	Tier 3	Required	Yes	High	\$20,000
			No	High	\$15,000
			No	Low	\$5,000

Denial of Service and Information Disclosure


The vulnerability should result in one of the following:

- Crash the host machine, resulting in a denial of service condition
- Cause a failure to start and stop VMs
- Gain sensitive information from the host machine or another guest

Vulnerability Type	Tier	Proof of concept	Report Quality	Payout range (USD)
DOS	Tier 1	Required	High	\$15,000
			Low	\$5,000
Info Disclosure	Tier 1	Required	High	\$25,000
			Low	\$5,000
	Tier 2	Required	High	\$15,000
			Low	\$5,000


- ✓ 1 x \$200,000
- ✓ 1 x \$15,000

How Much Are They Worth? (WIP Bounty)



- ✓ 2 x \$20,000
- ✓ 5 x \$20,000

CVE-2018-8439



CVE-2018-8439 | Windows Hyper-V Remote Code Execution Vulnerability Security Vulnerability

Published: 09/11/2018

MITRE CVE-2018-8439

A remote code execution vulnerability exists when Windows Hyper-V on a host server fails to properly validate input from an authenticated user on a guest operating system. To exploit the vulnerability, an attacker could run a specially crafted application on a guest operating system that could cause the Hyper-V host operating system to execute arbitrary code.

An attacker who successfully exploited the vulnerability could execute arbitrary code on the host operating system.

The security update addresses the vulnerability by correcting how Hyper-V validates guest operating system user input.

Exploitability Assessment

The following table provides an [exploitability assessment](#) for this vulnerability at the time of original publication.

Publicly Disclosed	Exploited	Latest Software Release	Older Software Release	Denial of Service
No	No	2 - Exploitation Less Likely	2 - Exploitation Less Likely	Not Applicable

On this page

[Executive Summary](#)

[Exploitability Assessment](#)

[Security Updates](#)

[Mitigations](#)

[Workarounds](#)

[FAQ](#)

Acknowledgements

Daniel King (@long123king), MSRC Microsoft ⁵⁷

CVE-2018-8439

HvPostMessage

- The hypervisor provides a simple inter-partition communication facility that allows one partition to send a parameterized message to another partition. (Because the message is sent asynchronously, it is said to be posted.) The destination partition may be notified of the arrival of this message through an SynIC(*Synthetic Interrupt Controller*) interrupt.

FROM: <https://docs.microsoft.com/en-us/virtualization/hyper-v-on-windows/reference/tlf>

Wrapper Interface

```
HV_STATUS  
HvPostMessage(  
    __in HV_CONNECTION_ID    ConnectionId,  
    __in HV_MESSAGE_TYPE     MessageType,  
    __in UINT32               PayloadSize,  
    __in_ecount(Payloadsize)  
        PCVOID Message  
);
```

Native Interface

HvPostMessage		
Call Code = 0x005C		
► Input Parameters		
0	ConnectionId (4 bytes)	Padding (4 bytes)
8	MessageType (4 bytes)	PayloadSize (4 bytes)
16	Message[0] (8 bytes)	
:	:	
248	Message[29] (8 bytes)	

CVE-2018-8439

- Reproduce Steps
 - I. Run Hyperseed inside guest VM to keep fuzzing HvPostMessage with ChannelMessageOpenChannel(0x05)/ChannelMessageCloseChannel(0x07) continuously.
 - II. Manually reset guest VM from host.
- Result
 - I. Host BSOD

CVE-2018-8439

```
0: kd> !analyze -v
*****
*           Bugcheck Analysis
*
*****  
  
DRIVER_IRQL_NOT_LESS_OR_EQUAL (d1)
An attempt was made to access a pageable (or completely invalid) address at an
interrupt request level (IRQL) that is too high. This is usually
caused by drivers using improper addresses.
If kernel debugger is available get stack backtrace.
Arguments:
Arg1: fffffba81d8001008, memory referenced
Arg2: 0000000000000002, IRQL
Arg3: 0000000000000001, value 0 = read operation, 1 = write operation
Arg4: fffff80fb4b51122, address which referenced memory
0: kd> .trap 0xfffff58dbc6ca7a0
NOTE: The trap frame does not contain all registers.
Some register values may be zeroed or incorrect.
rax=fffff8780eb36adc0 rbx=0000000000000000 rcx=0000000000000000
rdx=fffffba81d8001008 rsi=0000000000000000 rdi=0000000000000000
rip=fffff80fb4b51122 rsp=fffff58dbc6ca930 rbp=0000000080000005
r8=0000000000000001 r9=0000000000000000 r10=000000000000a000
r11=fffffba81d7a6a010 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0          nv up ei ng nz na pe nc
vmbusr!BusChSendInterrupt+0x42:
fffff80f`b4b51122 f00fab0a  lock bts dword ptr [rdx],ecx ds:fffffba81`d8001008=??
0: kd> kf
*** Stack trace for last set context - .thread/.cxr resets it
#  Memory  Child-SP      RetAddr      Call Site
00    fffff58d`bc6ca930  fffff80f`b5542e5c vmbusr!BusChSendInterrupt+0x42
01    30 fffff58d`bc6ca960  fffff80f`b55429a5 vmbkmclr!IncompletePacket+0x4ac
02    c0 fffff58d`bc6caa20  fffff80f`b72f112b vmbkmclr!VmbChannelPacketComplete+0x15
03    30 fffff58d`bc6caa50  fffff80f`b55413a2 storvsp+0x112b
04    60 fffff58d`bc6caaab0  fffff80f`b4b55680 vmbkmclr!InpProcessingWorkerRoutine+0x212
05    80 fffff58d`bc6cab30  fffff800`d28fdcae7 vmbusr!AwWorkerThread+0xb0
06    60 fffff58d`bc6cab90  fffff800`d29bbb86 nt!PspSystemThreadStartup+0x47
07    50 fffff58d`bc6cabe0  00000000`00000000 nt!KiStartSystemThread+0x16
0: kd> !pte rdx
VA fffffba81d8001008
PXE at FFFF98CC66331BA8    PPE at FFFF98CC66375038    PDE at FFFF98CC6EA07600
contains 0A0000000515D863  contains 0A0000000515E863  contains 0A0000023ACB5863
pfn 515d      ---DA--KWEV pfn 515e      ---DA--KWEV pfn 23acb5      ---DA--KWEV
                                         PTE at FFFF98DD40EC0008
                                         contains 001D845D00000000
                                         not valid
                                         Page has been freed
```

```
0: kd> dt vmbusr!_HV_MONITOR_PAGE -b
+0x000 TriggerState      : _HV_MONITOR_TRIGGER_STATE
+0x000 AsUINT32          : Uint4B
+0x000 GroupEnable       : Pos 0, 4 Bits
+0x000 RsvdZ             : Pos 4, 28 Bits
+0x004 RsvdZ1            : Uint4B
+0x008 TriggerGroup      : _HV_MONITOR_TRIGGER_GROUP
+0x000 AsUINT64          : Uint8B
+0x000 Pending            : Uint4B
+0x004 Armed              : Uint4B
+0x028 RsvdZ2            : Uint8B
+0x040 NextCheckTime     : Int4B
+0x240 Latency            : Uint2B
+0x340 RsvdZ3            : Uint8B
+0x440 Parameter          : _HV_MONITOR_PARAMETER
+0x000 ConnectionId      : _HV_CONNECTION_ID
+0x000 AsUINT32          : Uint4B
+0x000 Id                 : Pos 0, 24 Bits
+0x000 Reserved          : Pos 24, 8 Bits
+0x004 FlagNumber         : Uint2B
+0x006 RsvdZ              : Uint2B
+0x840 RsvdZ4            : UChar
```

CVE-2018-8439

- Monitored Notification Page
 - Allocated in vmbusr!ParentCreateMonitors
 - When initializing VMBus connection from the client
 - Freed in vmbusr!ParentDeleteMonitors
 - When resetting, destroying partition or failing VMBus connection
 - Cache a reference in vmbusr!PncEnableInterrupt
 - BusFdoOpenChannel, holds the reference even if the channel is not opened successfully
 - Clear the cached reference in vmbusr!PncDisableInterrupt
 - BusFdoCloseChannel, won't call this routine if it is not opened successfully
 - Accessing the cached reference after being freed leads to UAF

CVE-2018-8439 (Fix)

- Monitored Notification Page
 - Allocated in vmbusr!ParentCreateMonitors
 - When initializing VMBus connection
 - Freed in vmbusr!ParentDeleteMonitors
 - When resetting, destroying partition or failing VMBus connection
 - Cache a reference in vmbusr!PncEnableInterrupt
 - BusFdoOpenChannel, holds the reference even if the channel not opened successfully
 - Clear the cached reference in vmbusr!PncDisableInterrupt
 - BusFdoCloseChannel, won't be called if the channel is not opened successfully
 - **BusFdoOpenResult, clear the cached reference if the channel is not opened properly.**

Future



Future

- Keep Improving Format-Aware Fuzzing
- Adopt Coverage-Guided Fuzzing
 - Intel PT & AFL
- More Details on Findings After Being Patched
- Open-Sourcing
- More Targets

Acknowledgements

- Microsoft Security Response Center
- CDG Platform Security & Vulnerability Research
- Hyper-V Development Team

Appendix

Past Research

- [Security Assessment of Microsoft Hyper-V \(ERNW Newsletter 43 / May 2014, Felix Wilhelm, Mattias Luft\)](#)
- [Ring 0 to Ring -1 \(Syscan 2015 – Alex Ionescu\)](#)
- [The Battle of SKM and IUM \(Blackhat 2015 - Alex Ionescu\)](#)
- [Analysis of the Attack Surface of Windows 10 Virtualization-Base Security \(Blackhat 2016 - Rafal Wojtczuk\)](#)
- [VBS and VSM Internals \(Saar Amar\)](#)
- [A Dive in to Hyper-V Architecture and Vulnerabilities \(Blackhat 2018 - Joe Bialek & Nicolas Joly\)](#)
- [Hardening Hyper-V through offensive security research \(Blackhat 2018 – Jordan Rabet\)](#)
- [First steps in Hyper-V research \(Saar Amar\)](#)
- [Fuzzing para-virtualized devices in Hyper-V \(MSFT Virtualization Security Team\)](#)
- [Writing a Hyper-V “Bridge” for Fuzzing \(Alex Ionescu\)](#)

Useful Public Resources

- Hypervisor Top-Level Functional Specification
 - <https://docs.microsoft.com/en-us/virtualization/hyper-v-on-windows/reference/tlfs>
- Hyper-V Debugging Symbols
 - <https://blogs.technet.microsoft.com/virtualization/2018/04/25/hyper-v-symbols-for-debugging/>
- Linux Integration Service for Microsoft Hyper-V
 - <https://github.com/LIS>
- Unofficial hdk – Hyper-V Development Kit
 - <https://ionescu007.github.io/hdk/>
- Report quality definitions for Microsoft's Bug Bounty programs
 - <https://www.microsoft.com/en-us/msrc/bounty-example-report-submission>

Q&A



